

Comparing Deep Reinforcement Learning and Traditional Methods for Adaptive Portfolio Management

Blaine Swieder

Department of Computer Science
College of Engineering
Tennessee Technological University
Cookeville, Tennessee 38505
Email: bcswieder42@tntech.edu

Abstract—In this project, we will aim to provide a comparative analysis of Deep Reinforcement Learning algorithms for optimizing algorithmic trading strategies; more specifically, Proximal Policy Optimization, Advantage Actor-Critic, and Deep Q-Networks. Our principal goal is to identify which deep learning architecture delivers superior risk-adjusted returns when applied to dynamic financial markets. Moreover, to establish a performance benchmark, we will utilize traditional machine learning models such as gradient boosting machines and deep neural networks to provide further evaluation. We will utilize historical price data, technical indicators (that is, MACD, RSI), and macroeconomic variables, as well as interest fluctuations, to create a robust trading environment. The models will be trained and tested across multiple different market regimes, bullish, bearish, and sideways, to allow a comprehensive assessment of adaptability and robustness. In addition, we will measure performance using key financial metrics such as the Sharpe Ratio, Sortino Ratio, Maximum Drawdown, and annualized returns. In addition, this project will incorporate cost analysis to assess real-world trading viability and evaluate the impact of market volatility on algorithmic performance. The expected outcome of this project is to generate empirical insights into the comparative strengths and limitations of each algorithm, which will offer practical guidance for developing robust data-driven trading strategies in dynamic financial environments.

Index Terms—Algorithmic Trading, Machine Learning, Quantitative Finance, Risk Management

1. Problem Statement and Significance

1.1. Problem

When it comes to stock trading, modern financial markets can present unique challenges for algorithmic trading systems where the effective implementation of machine learning involves balancing three competing objectives:

- 1) Maximizing risk-adjusted returns (Sharpe Ratio > 1.5)
- 2) Maintaining sub-10 ms execution latency

- 3) Ensuring model stability across market regimes

The main objective of this project is to address the following question: *Which reinforcement learning algorithm, compared to traditional machine learning algorithms, provides the most optimal performance in different trading frequencies and market volatility regimes?* It is essential to address this question, as financial markets are inherently dynamic, thus showing sudden shifts in volatility, liquidity, and price behavior. These objectives are difficult to balance, since models optimized for high returns often suffer from latency issues, while those designed for low latency may compromise predictive accuracy. Algorithmic trading models that perform well under specific conditions may fail dramatically when exposed to different market environments will lead to significant financial losses. This will lose both the user and the investor's confidence in the system.

1.2. Significance

This study is of great practical importance for both institutional and retail traders, etc. Every year, institutional investors lose approximately \$2.1 billion annually due to suboptimal algorithm selection, which highlights the financial impact of poorly performing models [8]. In addition, around 68% of retail trading platforms lack robust model evaluation frameworks, which result in increased exposure to market risks and inconsistent trading outcomes. From a legal and regulatory perspective, compliance with policies such as SEC Rule 15c3-5 mandates the implementation of explainable trading strategies that ensure transparency and prevent the market from being manipulation.

The aim of this project is picking which machine learning algorithm—amongst Proximal Policy Optimization (PPO), Advantage Actor-Critic (A2C), Deep Q-Networks (DQN), Gradient Boosting Machines (GBM), and Deep Neural Networks (DNN) – consistently has superior performance across various financial conditions. Every one of these machine learning algorithms has its proven strengths. For example, deep reinforcement learning models such as Proximal Policy Optimization (PPO) and Advantage Actor-Critic (A2C) usually excel in dynamic environments

with continuous learning capabilities. Moreover, Gradient Boosting Machines (GBM) offer strong predictive power for structured data while Deep Neural Networks (DNN) capture complex nonlinear patterns. While prior literature has demonstrated the effectiveness of individual models, relatively few offers a comprehensive comparative analysis across multiple market conditions. If we evaluate these algorithms across diverse market regimes, such as bullish, bearish, and sideways; then, we will be able to identify not only top performer overall, but we also identify how each model effectiveness varies with market volatility, execution constraints, and risk profiles. By completing this study, we will be able to provide critical insights into algorithm selection for developing resilient, high-performance trading strategies for investors and users.

2. Background and Related Work

In the literature, deep reinforcement learning has emerged as a transformational approach in algorithmic trading due to its ability to be able to learn optimal trading strategies directly from dynamic market environments. Deep reinforcement learning methods, such as Proximal Policy Optimization (PPO), Advantage Actor-Critic (A2C), and Deep Q-Networks (DQN) have been proven to show significant promise in handling non-stationary data, adapting to changing market conditions, and maximizing long-term returns. As opposed to traditional machine learning models that rely heavily on historical data for predictive accuracy, deep reinforcement learning models continuously learn from the interactions with the market environment hence allowing them to adapt to real-time fluctuations [1], [2].

Moreover, the literature also reveals that Deep Reinforcement Learning outperforms conventional models in scenarios characterized by high volatility and complex market dynamics. For example, Zhang et al. [3] demonstrated the effectiveness of deep reinforcement learning algorithms in volatile markets, attributing their success to the ability of reinforcement learning to optimize decision-making under uncertainty. Similarly, Patel et al. [4] demonstrated that integrating deep neural networks with deep reinforcement learning frameworks enhances predictive accuracy, especially in high-frequency trading applications. In addition, traditional models such as Gradient Boosting Machines (GBM) and Deep Neural Networks (DNN), despite being robust in structured data analysis, often lacks the adaptability that is required in rapidly changing financial environments [7].

More recent work suggests that hybrid models that combine deep reinforcement learning with traditional techniques can mitigate challenges such as overfitting, computational inefficiency, and sensitivity to hyper parameters [6]. This research aims to build upon existing work by systematically comparing the performance of deep reinforcement learning algorithms with traditional models using different market regimes and by aiming to provide a comprehensive understanding of their strengths, limitations, as well as practical applications in algorithmic trading.

3. Data Sources and Methodology

For this project, we will utilize a dataset that is comprised of historical price data that is sourced from reputable financial databases like Bloomberg, Yahoo Finance, and Quandl. In addition, macroeconomic indicators, which include interest rate fluctuations and GDP growth rates that are obtained from reliable sources such as the Federal Reserve Economic Data (FRED) and World Bank datasets; moreover, technical indicators like Moving Average Convergence Divergence (MACD), Relative Strength Index (RSI), and Bollinger Bands are incorporated to capture market trends and volatility.

Our data preprocessing will involve normalization, feature engineering, and noise reduction techniques to enhance model performance. We could employ a rolling-window approach to simulate real-time trading conditions, to ensure robustness against data leakage and providing a realistic evaluation of model adaptability. The DRL algorithms are trained alongside traditional models such as GBM and DNNs to establish performance benchmarks, as done shown by Lee et al. and Prado [2], [7].

The model performance will be evaluated using key financial metrics, which includes the Sharpe Ratio, Sortino Ratio, Maximum Drawdown, and annualized returns. This methodological framework is intended to provide a rigorous, empirical basis to evaluate the adaptability, robustness, and practical viability of machine learning algorithms in diverse financial market conditions.

4. Conclusion

In conclusion, this project will aim to provide a comparative analysis of Deep Reinforcement Algorithms such as Proximal Policy Optimization (PPO), Advantage Actor Critic (A2C), and Deep Q-Networks (DQN) for optimizing algorithmic trading strategies. Our goal will be to identify which DRL architecture will provide the superior risk-adjusted returns in dynamic financial markets. To benchmark performance, we will use traditional machine learning models such as Gradient Boosting Machines (GBMs) and Deep Neural Networks (DNNs) as a comparison for the evaluation.

For our data, we will use historical price data from Bloomberg, Yahoo Finance, and Quandl, as well as technical indicators (MACD, RSI) and macroeconomic variables (i.e., interest rate fluctuations from FRED and World Bank). We will also test across diverse market regimes such as bullish, bearish, and sideways to determine and assess adaptability and robustness. Some key performance metrics such as the Sharpe Ratio, Sortino Ratio, Maximum Drawdown, and annualized returns.

Finally, the expected outcome will be to generate empirical insights into the advantages and disadvantages of each algorithm. We will aim to guide the development of robust, data-driven trading strategies that will adapt effectively to dynamic market conditions, which will contribute to more resilient and efficient trading systems.

References

- [1] M. Smith and L. Johnson, "Deep Reinforcement Learning Robots for Algorithmic Trading: Considering Stock Market Conditions and U.S. Interest Rates," *Journal of Financial Data Science*, vol. 3, no. 1, pp. 15–27, 2021.
- [2] R. Lee and S. Kim, "Algorithmic Trading on Financial Time Series Using Deep Reinforcement Learning," *IEEE Access*, vol. 8, pp. 212450–212460, 2020.
- [3] X. Zhang, Y. Liu, and Q. Wang, "A Novel Deep Reinforcement Learning-Based Automatic Stock Trading Method and a Case Study," *Applied Soft Computing*, vol. 94, pp. 106–119, 2020.
- [4] D. Patel and A. Banerjee, "High-Frequency Trading Strategy Based on Deep Neural Networks," *Quantitative Finance*, vol. 22, no. 4, pp. 567–578, 2022.
- [5] J. Thompson, "Deep Reinforcement Learning in Quantitative Algorithmic Trading: A Review," *Finance and AI Review*, vol. 5, no. 2, pp. 34–48, 2021.
- [6] K. Wang and H. Zhao, "Gradient Boosting Reinforcement Learning," *Proceedings of the International Conference on Machine Learning*, 2020.
- [7] M. Lopez de Prado, *Advances in Financial Machine Learning*. Hoboken, NJ: Wiley, 2018.
- [8] J.P. Morgan, "Annual Report on Institutional Trading Strategies," J.P. Morgan, 2023.
- [9] Bloomberg, "Bloomberg Markets - Financial Data & Analytics," *Bloomberg*, 2024. [Online]. Available: <https://www.bloomberg.com/>. [Accessed: Feb. 3, 2025].
- [10] Yahoo Finance, "Stock Market Data - Historical and Real-Time Prices," *Yahoo Finance*, 2024. [Online]. Available: <https://finance.yahoo.com/>. [Accessed: Feb. 3, 2025].
- [11] Quandl, "Financial and Economic Data API," *Quandl*, 2024. [Online]. Available: <https://www.quandl.com/>. [Accessed: Feb. 3, 2025].
- [12] Federal Reserve Bank of St. Louis, "FRED Economic Data," *Federal Reserve Economic Data (FRED)*, 2024. [Online]. Available: <https://fred.stlouisfed.org/>. [Accessed: Feb. 3, 2025].
- [13] The World Bank, "World Bank Open Data," *The World Bank*, 2024. [Online]. Available: <https://data.worldbank.org/>. [Accessed: Feb. 3, 2025].